

```

import numpy as np

class Softmax(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the Softmax classifier.
        Note that it has shape (C, D) where C is the number of
        classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims) * 0.0001

    def loss(self, X, y):
        """
        Calculates the softmax loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
            that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """
        # Initialize the loss to zero.
        loss = 0.0

        # ===== #
        # YOUR CODE HERE:
        # Calculate the normalized softmax loss. Store it as the variable loss.
        # (That is, calculate the sum of the losses of all the training
        # set margins, and then normalize the loss by the number of
        # training examples.)
        # ===== #
        weights = self.W
        losses = []

        for index in range(0, X.shape[0]):
            data = X[index]
            multiplication = weights.dot(data.T).T
            exponentials = np.exp(multiplication)
            exponentialSum = np.sum(exponentials)
            exponentials = exponentials / exponentialSum
            label = y[index]
            crossEntropy = -np.log(exponentials[label])
            losses.append(crossEntropy)
        loss = sum(losses) / len(losses)

        # ===== #
        # END YOUR CODE HERE

```

```

# ===== #

return loss

def loss_and_grad(self, X, y):
    """
    Same as self.loss(X, y), except that it also returns the gradient.

    Output: grad -- a matrix of the same dimensions as W containing
    the gradient of the loss with respect to W.
    """

    # Initialize the loss and gradient to zero.
    loss = 0.0
    grad = np.zeros_like(self.W)

    # ===== #
    # YOUR CODE HERE:
    # Calculate the softmax loss and the gradient. Store the gradient
    # as the variable grad.
    # ===== #
    weights = self.W
    losses=[]
    for index in range(0,X.shape[0]):
        data = X[index]
        multiplication = weights.dot(data.T).T
        exponentials = np.exp(multiplication)
        exponentialSum = np.sum(exponentials)
        exponentials=exponentials/exponentialSum
        label = y[index]
        crossEntropy = -np.log(exponentials[label])
        losses.append(crossEntropy)
        yHot = np.zeros([1,np.max(y)+1])
        yHot[0,label]=1
        data = np.reshape(data,[1,data.shape[0]])
        exponentials = np.reshape(exponentials,[1,exponentials.shape[0]])
        derivative = np.dot(data.T, (exponentials - yHot))
        grad=grad+derivative.T
    loss = sum(losses)/len(losses)
    grad = grad/[X.shape[0]]

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical
    in these dimensions.
    """

    for i in np.arange(num_checks):
        ix = tuple([np.random.randint(m) for m in self.W.shape])

        oldval = self.W[ix]

```

```

self.W[ix] = oldval + h # increment by h
fxph = self.loss(X, y)
self.W[ix] = oldval - h # decrement by h
fxmh = self.loss(X,y) # evaluate f(x - h)
self.W[ix] = oldval # reset

grad_numerical = (fxph - fxmh) / (2 * h)
grad_analytic = your_grad[ix]
rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) + abs(grad_analytic))
print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical, grad_analytic, rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and outputs as loss_and_grad.
    """
    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero

    # ===== #
    # YOUR CODE HERE:
    # Calculate the softmax loss and gradient WITHOUT any for loops.
    # ===== #
    weights = self.W
    multiplication = weights.dot(X.T).T
    multiplication = (multiplication.T - np.amax(multiplication,axis = 1)).T
    soft = np.exp(multiplication)
    sums = np.sum(soft,axis=1)
    probs = (soft.T / sums).T
    predsForClass = probs[np.arange(y.size),y]
    loss = -np.log(predsForClass+1e-10) #To avoid log(0)
    loss = np.mean(loss)
    yHot = np.zeros([y.shape[0],9+1])
    yHot[np.arange(y.size),y] = 1
    grad = (1/X.shape[0])*np.dot(X.T, (probs - yHot)).T

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there are N
        training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i] = c
        means that X[i] has label 0 <= c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each step.
    - verbose: (boolean) If true, print progress during optimization.
    """

```

Outputs:

A list containing the value of the loss function at each training iteration.

"""

num_train, dim = X.shape

num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes

self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of self.W

Run stochastic gradient descent to optimize W

loss_history = []

for it in np.arange(num_iters):

 X_batch = None

 y_batch = None

 # ===== #

 # YOUR CODE HERE:

 # Sample batch_size elements from the training data for use in
 # gradient descent. After sampling,

 # - X_batch should have shape: (dim, batch_size)

 # - y_batch should have shape: (batch_size,)

 # The indices should be randomly generated to reduce correlations
 # in the dataset. Use np.random.choice. It's okay to sample with
 # replacement.

 # ===== #

 indices = np.random.choice(np.arange(num_train), batch_size)

 XBatch = X[indices]

 yBatch = y[indices]

 # ===== #

 # END YOUR CODE HERE

 # ===== #

 # evaluate loss and gradient

 loss, grad = self.fast_loss_and_grad(XBatch, yBatch)

 loss_history.append(loss)

 # ===== #

 # YOUR CODE HERE:

 # Update the parameters, self.W, with a gradient step

 # ===== #

 self.W = self.W - learning_rate*grad

 # ===== #

 # END YOUR CODE HERE

 # ===== #

 if verbose and it % 100 == 0:

 print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

return loss_history

def predict(self, X):

"""

Inputs:

- X: N x D array of training data. Each row is a D-dimensional point.

Returns:

- `y_pred`: Predicted labels for the data in `X`. `y_pred` is a 1-dimensional array of length `N`, and each element is an integer giving the predicted class.

```
"""
y_pred = np.zeros(X.shape[1])
# ===== #
# YOUR CODE HERE:
#   Predict the labels given the training data.
# ===== #
weights = self.W
multiplication = weights.dot(X.T).T
soft = np.exp(multiplication)
y_pred = np.argmax(soft, axis=1)
# ===== #
# END YOUR CODE HERE
# ===== #

return y_pred
```