

```

import numpy as np

from .layers import *
from .layer_utils import *

class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network with ReLU nonlinearity and
    softmax loss that uses a modular layer design. We assume an input dimension
    of D, a hidden dimension of H, and perform classification over C classes.

    The architecture should be affine - relu - affine - softmax.

    Note that this class does not implement gradient descent; instead, it
    will interact with a separate Solver object that is responsible for running
    optimization.

    The learnable parameters of the model are stored in the dictionary
    self.params that maps parameter names to numpy arrays.
    """

    def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
                  dropout=0, weight_scale=1e-3, reg=0.0):
        """
        Initialize a new network.

        Inputs:
        - input_dim: An integer giving the size of the input
        - hidden_dims: An integer giving the size of the hidden layer
        - num_classes: An integer giving the number of classes to classify
        - dropout: Scalar between 0 and 1 giving dropout strength.
        - weight_scale: Scalar giving the standard deviation for random
          initialization of the weights.
        - reg: Scalar giving L2 regularization strength.
        """
        self.params = {}
        self.reg = reg

        # ===== #
        # YOUR CODE HERE:
        # Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
        # self.params['W2'], self.params['b1'] and self.params['b2']. The
        # biases are initialized to zero and the weights are initialized
        # so that each parameter has mean 0 and standard deviation weight_scale.
        # The dimensions of W1 should be (input_dim, hidden_dim) and the
        # dimensions of W2 should be (hidden_dims, num_classes)
        # ===== #

        self.params['W1'] = np.random.normal(loc=0.0, scale=weight_scale, size = (input_dim, hidden_dims))
        self.params['b1'] = np.zeros(hidden_dims)
        self.params['W2'] = np.random.normal(loc=0.0, scale=weight_scale, size = (hidden_dims, num_classes))
        self.params['b2'] = np.zeros(num_classes)

        # ===== #

```

```

# END YOUR CODE HERE
# ===== #

def loss(self, X, y=None):
    """
    Compute loss and gradient for a minibatch of data.

    Inputs:
    - X: Array of input data of shape (N, d_1, ..., d_k)
    - y: Array of labels, of shape (N,). y[i] gives the label for X[i].

    Returns:
    If y is None, then run a test-time forward pass of the model and return:
    - scores: Array of shape (N, C) giving classification scores, where
      scores[i, c] is the classification score for X[i] and class c.

    If y is not None, then run a training-time forward and backward pass and
    return a tuple of:
    - loss: Scalar value giving the loss
    - grads: Dictionary with the same keys as self.params, mapping parameter
      names to gradients of the loss with respect to those parameters.
    """
    scores = None

    # ===== #
    # YOUR CODE HERE:
    #   Implement the forward pass of the two-layer neural network. Store
    #   the class scores as the variable 'scores'. Be sure to use the layers
    #   you prior implemented.
    # ===== #
    W1 = self.params['W1']
    b1 = self.params['b1']
    W2 = self.params['W2']
    b2 = self.params['b2']

    a, fc_cache = affine_forward(X, W1, b1)
    out, relu_cache = relu_forward(a)
    cache_hidden = (fc_cache, relu_cache)
    scores, cache_z = affine_forward(out, W2, b2)

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    # If y is None then we are in test mode so just return scores
    if y is None:
        return scores

    loss, grads = 0, {}

    # ===== #
    # YOUR CODE HERE:
    #   Implement the backward pass of the two-layer neural net. Store
    #   the loss as the variable 'loss' and store the gradients in the
    #   'grads' dictionary. For the grads dictionary, grads['W1'] holds
    #   the gradient for W1, grads['b1'] holds the gradient for b1, etc.
    #   i.e., grads[k] holds the gradient for self.params[k].

```

```

#
# Add L2 regularization, where there is an added cost 0.5*self.reg*W^2
# for each W. Be sure to include the 0.5 multiplying factor to
# match our implementation.
#
# And be sure to use the layers you prior implemented.
# ===== #

loss, dz = softmax_loss(scores, y)
loss = loss + 0.5*self.reg*(np.sum(W1**2) + np.sum(W2**2))

dhidden, dw2, db2 = affine_backward(dz, cache_z)
fc_cache, relu_cache = cache_hidden
da = relu_backward(dhidden, relu_cache)
dx, dw1, db1 = affine_backward(da, fc_cache)

grads['W1'] = dw1 + self.reg * W1
grads['b1'] = db1
grads['W2'] = dw2 + self.reg * W2
grads['b2'] = db2
# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

```

```
class FullyConnectedNet(object):
```

```

"""
A fully-connected neural network with an arbitrary number of hidden layers,
ReLU nonlinearities, and a softmax loss function. This will also implement
dropout and batch normalization as options. For a network with L layers,
the architecture will be

```

```
{affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
```

```
where batch normalization and dropout are optional, and the {...} block is
repeated L - 1 times.
```

```

Similar to the TwoLayerNet above, learnable parameters are stored in the
self.params dictionary and will be learned using the Solver class.
"""

```

```

def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
              dropout=0, use_batchnorm=False, reg=0.0,
              weight_scale=1e-2, dtype=np.float32, seed=None):

```

```

    """

```

```
    Initialize a new FullyConnectedNet.
```

```
    Inputs:
```

- hidden_dims: A list of integers giving the size of each hidden layer.
- input_dim: An integer giving the size of the input.
- num_classes: An integer giving the number of classes to classify.
- dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then the network should not use dropout at all.
- use_batchnorm: Whether or not the network should use batch normalization.
- reg: Scalar giving L2 regularization strength.

- `weight_scale`: Scalar giving the standard deviation for random initialization of the weights.
- `dtype`: A numpy datatype object; all computations will be performed using this datatype. `float32` is faster but less accurate, so you should use `float64` for numeric gradient checking.
- `seed`: If not `None`, then pass this random seed to the dropout layers. This will make the dropout layers deterministic so we can gradient check the model.

"""

```
self.use_batchnorm = use_batchnorm
self.use_dropout = dropout > 0
self.reg = reg
self.num_layers = 1 + len(hidden_dims)
self.dtype = dtype
self.params = {}
```

=====

YOUR CODE HERE:

Initialize all parameters of the network in the `self.params` dictionary.
 # The weights and biases of layer 1 are `W1` and `b1`; and in general the
 # weights and biases of layer `i` are `Wi` and `bi`. The
 # biases are initialized to zero and the weights are initialized
 # so that each parameter has mean 0 and standard deviation `weight_scale`.

=====

```
for i in range(0, self.num_layers):
```

```
    name_W = 'W'+str(i+1)
    name_b = 'b'+str(i+1)
    if i == 0: #First
        self.params[name_W] = np.random.normal(loc=0.0, scale=weight_scale, size = (input
        self.params[name_b] = np.zeros(hidden_dims[i])
    elif i == self.num_layers-1: #Last
        self.params[name_W] = np.random.normal(loc=0.0, scale=weight_scale, size = (hidde
        self.params[name_b] = np.zeros(num_classes)
    else: #Between
        self.params[name_W] = np.random.normal(loc=0.0, scale=weight_scale, size = (hidde
        self.params[name_b] = np.zeros(hidden_dims[i])
```

=====

END YOUR CODE HERE

=====

When using dropout we need to pass a `dropout_param` dictionary to each
 # dropout layer so that the layer knows the dropout probability and the mode
 # (train / test). You can pass the same `dropout_param` to each dropout layer.

```
self.dropout_param = {}
if self.use_dropout:
    self.dropout_param = {'mode': 'train', 'p': dropout}
if seed is not None:
    self.dropout_param['seed'] = seed
```

With batch normalization we need to keep track of running means and
 # variances, so we need to pass a special `bn_param` object to each batch
 # normalization layer. You should pass `self.bn_params[0]` to the forward pass
 # of the first batch normalization layer, `self.bn_params[1]` to the forward
 # pass of the second batch normalization layer, etc.

```

self.bn_params = []
if self.use_batchnorm:
    self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]

# Cast all parameters to the correct datatype
for k, v in self.params.items():
    self.params[k] = v.astype(dtype)

def loss(self, X, y=None):
    """
    Compute loss and gradient for the fully-connected net.

    Input / output: Same as TwoLayerNet above.
    """
    X = X.astype(self.dtype)
    mode = 'test' if y is None else 'train'

    # Set train/test mode for batchnorm params and dropout param since they
    # behave differently during training and testing.
    if self.dropout_param is not None:
        self.dropout_param['mode'] = mode
    if self.use_batchnorm:
        for bn_param in self.bn_params:
            bn_param['mode'] = mode

    scores = None

    # ===== #
    # YOUR CODE HERE:
    # Implement the forward pass of the FC net and store the output
    # scores as the variable "scores".
    # ===== #

    H = []
    cache_h = []
    for i in np.arange(0, self.num_layers):
        name_W = 'W'+str(i+1)
        name_b = 'b'+str(i+1)

        if i == 0: #First
            a, fc_cache = affine_forward(X, self.params[name_W], self.params[name_b])
            out, relu_cache = relu_forward(a)
            cH = (fc_cache, relu_cache)
            H.append(out)
            cache_h.append(cH)
        elif i == self.num_layers-1: #Last
            scores = affine_forward(H[i-1], self.params[name_W], self.params[name_b])[0]
            cache_h.append(affine_forward(H[i-1], self.params[name_W], self.params[name_b]))
        else: #Between
            a, fc_cache = affine_forward(H[i-1], self.params[name_W], self.params[name_b])
            out, relu_cache = relu_forward(a)
            cH = (fc_cache, relu_cache)
            H.append(out)
            cache_h.append(cH)

    # ===== #

```

```

# END YOUR CODE HERE
# ===== #

# If test mode return early
if mode == 'test':
    return scores

loss, grads = 0.0, {}
# ===== #
# YOUR CODE HERE:
# Implement the backwards pass of the FC net and store the gradients
# in the grads dict, so that grads[k] is the gradient of self.params[k]
# Be sure your L2 regularization includes a 0.5 factor.
# ===== #

loss, dz = softmax_loss(scores, y)
for i in range(self.num_layers, 0, -1):
    name_W = 'W'+str(i)
    name_b = 'b'+str(i)
    loss = loss + (0.5 * self.reg * np.sum(self.params[name_W]*self.params[name_W]))

    if i == self.num_layers:
        dh1, grads[name_W], grads[name_b] = affine_backward(dz, cache_h[self.num_layers])
    else:
        fc_cache, relu_cache = cache_h[i-1]
        da = relu_backward(dh1, relu_cache)
        dh1, grads[name_W], grads[name_b] = affine_backward(da, fc_cache)

    grads[name_W] = grads[name_W] + self.reg * self.params[name_W]

# ===== #
# END YOUR CODE HERE
# ===== #
return loss, grads

```