

```
import numpy as np
import matplotlib.pyplot as plt
```

```
class TwoLayerNet(object):
```

```
    """
```

A two-layer fully-connected neural network. The net has an input dimension of N , a hidden layer dimension of H , and performs classification over C classes. We train the network with a softmax loss function and L2 regularization on the weight matrices. The network uses a ReLU nonlinearity after the first fully connected layer.

In other words, the network has the following architecture:

input - fully connected layer - ReLU - fully connected layer - softmax

The outputs of the second fully-connected layer are the scores for each class.

```
    """
```

```
def __init__(self, input_size, hidden_size, output_size, std=1e-4):
```

```
    """
```

Initialize the model. Weights are initialized to small random values and biases are initialized to zero. Weights and biases are stored in the variable `self.params`, which is a dictionary with the following keys:

`W1`: First layer weights; has shape (H, D)
`b1`: First layer biases; has shape $(H,)$
`W2`: Second layer weights; has shape (C, H)
`b2`: Second layer biases; has shape $(C,)$

Inputs:

- `input_size`: The dimension D of the input data.
- `hidden_size`: The number of neurons H in the hidden layer.
- `output_size`: The number of classes C .

```
    """
```

```
self.params = {}
self.params['W1'] = std * np.random.randn(hidden_size, input_size)
self.params['b1'] = np.zeros(hidden_size)
self.params['W2'] = std * np.random.randn(output_size, hidden_size)
self.params['b2'] = np.zeros(output_size)
```

```
def loss(self, X, y=None, reg=0.0):
```

```
    """
```

Compute the loss and gradients for a two layer fully connected neural network.

Inputs:

- `X`: Input data of shape (N, D) . Each `X[i]` is a training sample.
- `y`: Vector of training labels. `y[i]` is the label for `X[i]`, and each `y[i]` is an integer in the range $0 \leq y[i] < C$. This parameter is optional; if it is not passed then we only return scores, and if it is passed then we instead return the loss and gradients.
- `reg`: Regularization strength.

Returns:

If `y` is `None`, return a matrix scores of shape (N, C) where `scores[i, c]` is

the score for class `c` on input `X[i]`.

If `y` is not `None`, instead return a tuple of:

- `loss`: Loss (data loss and regularization loss) for this batch of training samples.
- `grads`: Dictionary mapping parameter names to gradients of those parameters with respect to the loss function; has the same keys as `self.params`.

```
"""
# Unpack variables from the params dictionary
W1, b1 = self.params['W1'], self.params['b1']
W2, b2 = self.params['W2'], self.params['b2']
N, D = X.shape

# Compute the forward pass
scores = None

# ===== #
# YOUR CODE HERE:
# Calculate the output scores of the neural network. The result
# should be (N, C). As stated in the description for this class,
# there should not be a ReLU layer after the second FC layer.
# The output of the second FC layer is the output scores. Do not
# use a for loop in your implementation.
# ===== #
#b1 = b1.reshape([len(b1),1])
#b2 = b2.reshape([len(b2),1])
layer1Out = W1.dot(X.T) + b1[:,None]
ReluOut = np.maximum(0,layer1Out)
layer2Out = W2.dot(ReluOut) + b2[:,None]
scores = layer2Out.T

# ===== #
# END YOUR CODE HERE
# ===== #

# If the targets are not given then jump out, we're done
if y is None:
    return scores

# Compute the loss
loss = None

# ===== #
# YOUR CODE HERE:
# Calculate the loss of the neural network. This includes the
# softmax loss and the L2 regularization for W1 and W2. Store the
# total loss in the variable loss. Multiply the regularization
# loss by 0.5 (in addition to the factor reg).
# ===== #

# scores is num_examples by num_classes
soft = np.exp(scores)
sums = np.sum(soft,axis=1)
probs = soft / sums[:,None]
predsForClass = probs[np.arange(y.shape[0]),y]
SoftmaxLoss = np.mean(-np.log(predsForClass))
```

```

l2regularization = np.sum(W1**2) + np.sum(W2**2)
l2regularization = 0.5*reg*l2regularization
loss=SoftmaxLoss+l2regularization
# ===== #
# END YOUR CODE HERE
# ===== #

grads = {}

# ===== #
# YOUR CODE HERE:
# Implement the backward pass. Compute the derivatives of the
# weights and the biases. Store the results in the grads
# dictionary. e.g., grads['W1'] should store the gradient for
# W1, and be of the same size as W1.
# ===== #
#A1 = W1X+B1
#A2 = RELU(A1)
#A3 = W2A2+B2
#A4 = SOFTMAX(A3)
#DL/DA3 = PREDICTIONS-LABELS ONEHOT
#DA3/DW2 = A2
#DA3/DA2 = W2
#DA2/DA1 = 0 OR 1
#DA1/DW1 = A1
grad = probs.copy()
grad[np.arange(y.shape[0]),y] -= 1
dLA3=grad/X.shape[0]
dA3W2=ReluOut.copy()
b2grad = np.sum(dLA3,axis=0)
w2grad = np.dot(dLA3.T, dA3W2.T) + reg*W2
dA3A2 = W2
dA2dA1 = layer10out.copy()
dA2dA1[dA2dA1<0]=0
dA2dA1[dA2dA1>0]=1
dA1dW1 = X.copy()

kronecker = ((dLA3.dot(dA3A2)).T*dA2dA1)
b1grad = np.sum(kronecker.T,axis=0)
w1grad = (kronecker.dot(dA1dW1)) + reg*W1

grads["W1"]=w1grad
grads["b1"]=b1grad
grads["W2"]=w2grad
grads["b2"]=b2grad

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

def train(self, X, y, X_val, y_val,
        learning_rate=1e-3, learning_rate_decay=0.95,
        reg=1e-5, num_iters=100,
        batch_size=200, verbose=False):
    """

```

Train this neural network using stochastic gradient descent.

Inputs:

- X: A numpy array of shape (N, D) giving training data.
- y: A numpy array of shape (N,) giving training labels; $y[i] = c$ means that $X[i]$ has label c , where $0 \leq c < C$.
- X_val: A numpy array of shape (N_val, D) giving validation data.
- y_val: A numpy array of shape (N_val,) giving validation labels.
- learning_rate: Scalar giving learning rate for optimization.
- learning_rate_decay: Scalar giving factor used to decay the learning rate after each epoch.
- reg: Scalar giving regularization strength.
- num_iters: Number of steps to take when optimizing.
- batch_size: Number of training examples to use per step.
- verbose: boolean; if true print progress during optimization.

```
num_train = X.shape[0]
iterations_per_epoch = max(num_train / batch_size, 1)

# Use SGD to optimize the parameters in self.model
loss_history = []
train_acc_history = []
val_acc_history = []

for it in np.arange(num_iters):
    X_batch = None
    y_batch = None

    # ===== #
    # YOUR CODE HERE:
    #   Create a minibatch by sampling batch_size samples randomly.
    # ===== #
    indices = np.random.choice(np.arange(num_train), batch_size)
    X_batch = X[indices]
    y_batch = y[indices]
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    # Compute loss and gradients using the current minibatch
    loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
    loss_history.append(loss)

    # ===== #
    # YOUR CODE HERE:
    #   Perform a gradient descent step using the minibatch to update
    #   all parameters (i.e., W1, W2, b1, and b2).
    # ===== #

    self.params['W1'] = self.params['W1'] - grads['W1'] * learning_rate
    self.params['b1'] = self.params['b1'] - grads['b1'] * learning_rate
    self.params['W2'] = self.params['W2'] - grads['W2'] * learning_rate
    self.params['b2'] = self.params['b2'] - grads['b2'] * learning_rate

    # ===== #
    # END YOUR CODE HERE
    # ===== #
```

```

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

# Every epoch, check train and val accuracy and decay learning rate.
if it % iterations_per_epoch == 0:
    # Check accuracy
    train_acc = (self.predict(X_batch) == y_batch).mean()
    val_acc = (self.predict(X_val) == y_val).mean()
    train_acc_history.append(train_acc)
    val_acc_history.append(val_acc)

    # Decay learning rate
    learning_rate *= learning_rate_decay

return {
    'loss_history': loss_history,
    'train_acc_history': train_acc_history,
    'val_acc_history': val_acc_history,
}

def predict(self, X):
    """
    Use the trained weights of this two-layer network to predict labels for
    data points. For each data point we predict scores for each of the C
    classes, and assign each data point to the class with the highest score.

    Inputs:
    - X: A numpy array of shape (N, D) giving N D-dimensional data points to
        classify.

    Returns:
    - y_pred: A numpy array of shape (N,) giving predicted labels for each of
        the elements of X. For all i, y_pred[i] = c means that X[i] is predicted
        to have class c, where 0 <= c < C.
    """
    y_pred = None

    # ===== #
    # YOUR CODE HERE:
    # Predict the class given the input data.
    # ===== #
    layer1out = self.params['W1'].dot(X.T) + self.params['b1'][:,None]
    ReluOut = np.maximum(0, layer1out)
    layer2out = self.params['W2'].dot(ReluOut) + self.params['b2'][:,None]
    y_pred = np.argmax(layer2out, axis = 0)

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return y_pred

```