

ECE C147/247 HW4 Q3: Dropout

In this notebook, you will implement dropout. Then we will ask you to train a network with batchnorm and dropout, and achieve over 55% accuracy on CIFAR-10.

`utils` has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, and their layer structure. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`.

```
In [1]: 1  ## Import and setups
2
3  import time
4  import numpy as np
5  import matplotlib.pyplot as plt
6  from nndl.fc_net import *
7  from nndl.layers import *
8  from utils.data_utils import get_CIFAR10_data
9  from utils.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
10 from utils.solver import Solver
11
12 %matplotlib inline
13 plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
14 plt.rcParams['image.interpolation'] = 'nearest'
15 plt.rcParams['image.cmap'] = 'gray'
16
17 # for auto-reloading external modules
18 # see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
19 %load_ext autoreload
20 %autoreload 2
21
22 def rel_error(x, y):
23     """ returns relative error """
24     return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

```
In [2]: 1  # Load the (preprocessed) CIFAR10 data.
2
3  data = get_CIFAR10_data()
4  for k in data.keys():
5      print('{}: {}'.format(k, data[k].shape))
```

```
X_train: (49000, 3, 32, 32)
y_train: (49000,)
X_val: (1000, 3, 32, 32)
y_val: (1000,)
X_test: (1000, 3, 32, 32)
y_test: (1000,)
```

Dropout forward pass

Implement the training and test time dropout forward pass, `dropout_forward`, in `nndl/layers.py`. After that, test your implementation by running the following cell.

```
In [3]: 1 x = np.random.randn(500, 500) + 10
2
3 for p in [0.3, 0.6, 0.75]:
4     out, _ = dropout_forward(x, {'mode': 'train', 'p': p})
5     out_test, _ = dropout_forward(x, {'mode': 'test', 'p': p})
6
7     print('Running tests with p = ', p)
8     print('Mean of input: ', x.mean())
9     print('Mean of train-time output: ', out.mean())
10    print('Mean of test-time output: ', out_test.mean())
11    print('Fraction of train-time output set to zero: ', (out == 0).mean())
12    print('Fraction of test-time output set to zero: ', (out_test == 0).mean())
```

```
Running tests with p = 0.3
Mean of input: 10.000869762098812
Mean of train-time output: 10.008144829587012
Mean of test-time output: 10.000869762098812
Fraction of train-time output set to zero: 0.699764
Fraction of test-time output set to zero: 0.0
Running tests with p = 0.6
Mean of input: 10.000869762098812
Mean of train-time output: 10.02942410584371
Mean of test-time output: 10.000869762098812
Fraction of train-time output set to zero: 0.398184
Fraction of test-time output set to zero: 0.0
Running tests with p = 0.75
Mean of input: 10.000869762098812
Mean of train-time output: 10.010225631186525
Mean of test-time output: 10.000869762098812
Fraction of train-time output set to zero: 0.249352
Fraction of test-time output set to zero: 0.0
```

Dropout backward pass

Implement the backward pass, `dropout_backward`, in `nndl/layer.py`. After that, test your gradients by running the following cell:

```
In [3]: 1 x = np.random.randn(10, 10) + 10
2 dout = np.random.randn(*x.shape)
3
4 dropout_param = {'mode': 'train', 'p': 0.8, 'seed': 123}
5 out, cache = dropout_forward(x, dropout_param)
6 dx = dropout_backward(dout, cache)
7 dx_num = eval_numerical_gradient_array(lambda xx: dropout_forward(xx, dropout_param)[0], x, dout)
8
9 print('dx relative error: ', rel_error(dx, dx_num))
```

```
dx relative error: 5.4456097039628486e-11
```

Implement a fully connected neural network with dropout layers

Modify the `FullyConnectedNet()` class in `nndl/fc_net.py` to incorporate dropout. A dropout layer should be incorporated after every ReLU layer. Concretely, there shouldn't be a dropout at the output layer since there is no ReLU at the output layer. You will need to modify the class in the following areas:

(1) In the forward pass, you will need to incorporate a dropout layer after every relu layer.

(2) In the backward pass, you will need to incorporate a dropout backward pass layer.

Check your implementation by running the following code. Our W1 gradient relative error is on the order of $1e-6$ (the largest of all the relative errors).

```
In [4]: 1 N, D, H1, H2, C = 2, 15, 20, 30, 10
2 X = np.random.randn(N, D)
3 y = np.random.randint(C, size=(N,))
4
5 for dropout in [0, 0.25, 0.5]:
6     print('Running check with dropout = ', dropout)
7     model = FullyConnectedNet([H1, H2], input_dim=D, num_classes=C,
8                               weight_scale=5e-2, dtype=np.float64,
9                               dropout=dropout, seed=123)
10
11     loss, grads = model.loss(X, y)
12     print('Initial loss: ', loss)
13
14     for name in sorted(grads):
15         f = lambda _: model.loss(X, y)[0]
16         grad_num = eval_numerical_gradient(f, model.params[name], verbose=False, h=1e-5)
17         print('{} relative error: {}'.format(name, rel_error(grad_num, grads[name])))
18     print('\n')
```

```
Running check with dropout = 0
Initial loss: 2.3051948273987857
W1 relative error: 2.5272575344376073e-07
W2 relative error: 1.5034484929313676e-05
W3 relative error: 2.753446833630168e-07
b1 relative error: 2.936957476400148e-06
b2 relative error: 5.051339805546953e-08
b3 relative error: 1.1740467838205477e-10
```

```
Running check with dropout = 0.25
Initial loss: 2.3126468345657742
W1 relative error: 1.483854795975875e-08
W2 relative error: 2.3427832149940254e-10
W3 relative error: 3.564454999162522e-08
b1 relative error: 1.5292167232408546e-09
b2 relative error: 1.842268868410678e-10
b3 relative error: 1.4026015558098908e-10
```

```
Running check with dropout = 0.5
Initial loss: 2.302437587710995
W1 relative error: 4.553387957138422e-08
W2 relative error: 2.974218050584597e-08
W3 relative error: 4.3413247403122424e-07
b1 relative error: 1.872462967441693e-08
b2 relative error: 5.045591219274328e-09
b3 relative error: 8.009887154529434e-11
```

Dropout as a regularizer

In class, we claimed that dropout acts as a regularizer by effectively bagging. To check this, we will train two small networks, one with dropout and one without dropout.

```

In [7]: 1 # Train two identical nets, one with dropout and one without
2
3 num_train = 500
4 small_data = {
5     'X_train': data['X_train'][:num_train],
6     'y_train': data['y_train'][:num_train],
7     'X_val': data['X_val'],
8     'y_val': data['y_val'],
9 }
10
11 solvers = {}
12 dropout_choices = [0, 0.6]
13 for dropout in dropout_choices:
14     model = FullyConnectedNet([100, 100, 100], dropout=dropout)
15
16     solver = Solver(model, small_data,
17                     num_epochs=25, batch_size=100,
18                     update_rule='adam',
19                     optim_config={
20                         'learning_rate': 5e-4,
21                     },
22                     verbose=True, print_every=100)
23     solver.train()
24     solvers[dropout] = solver

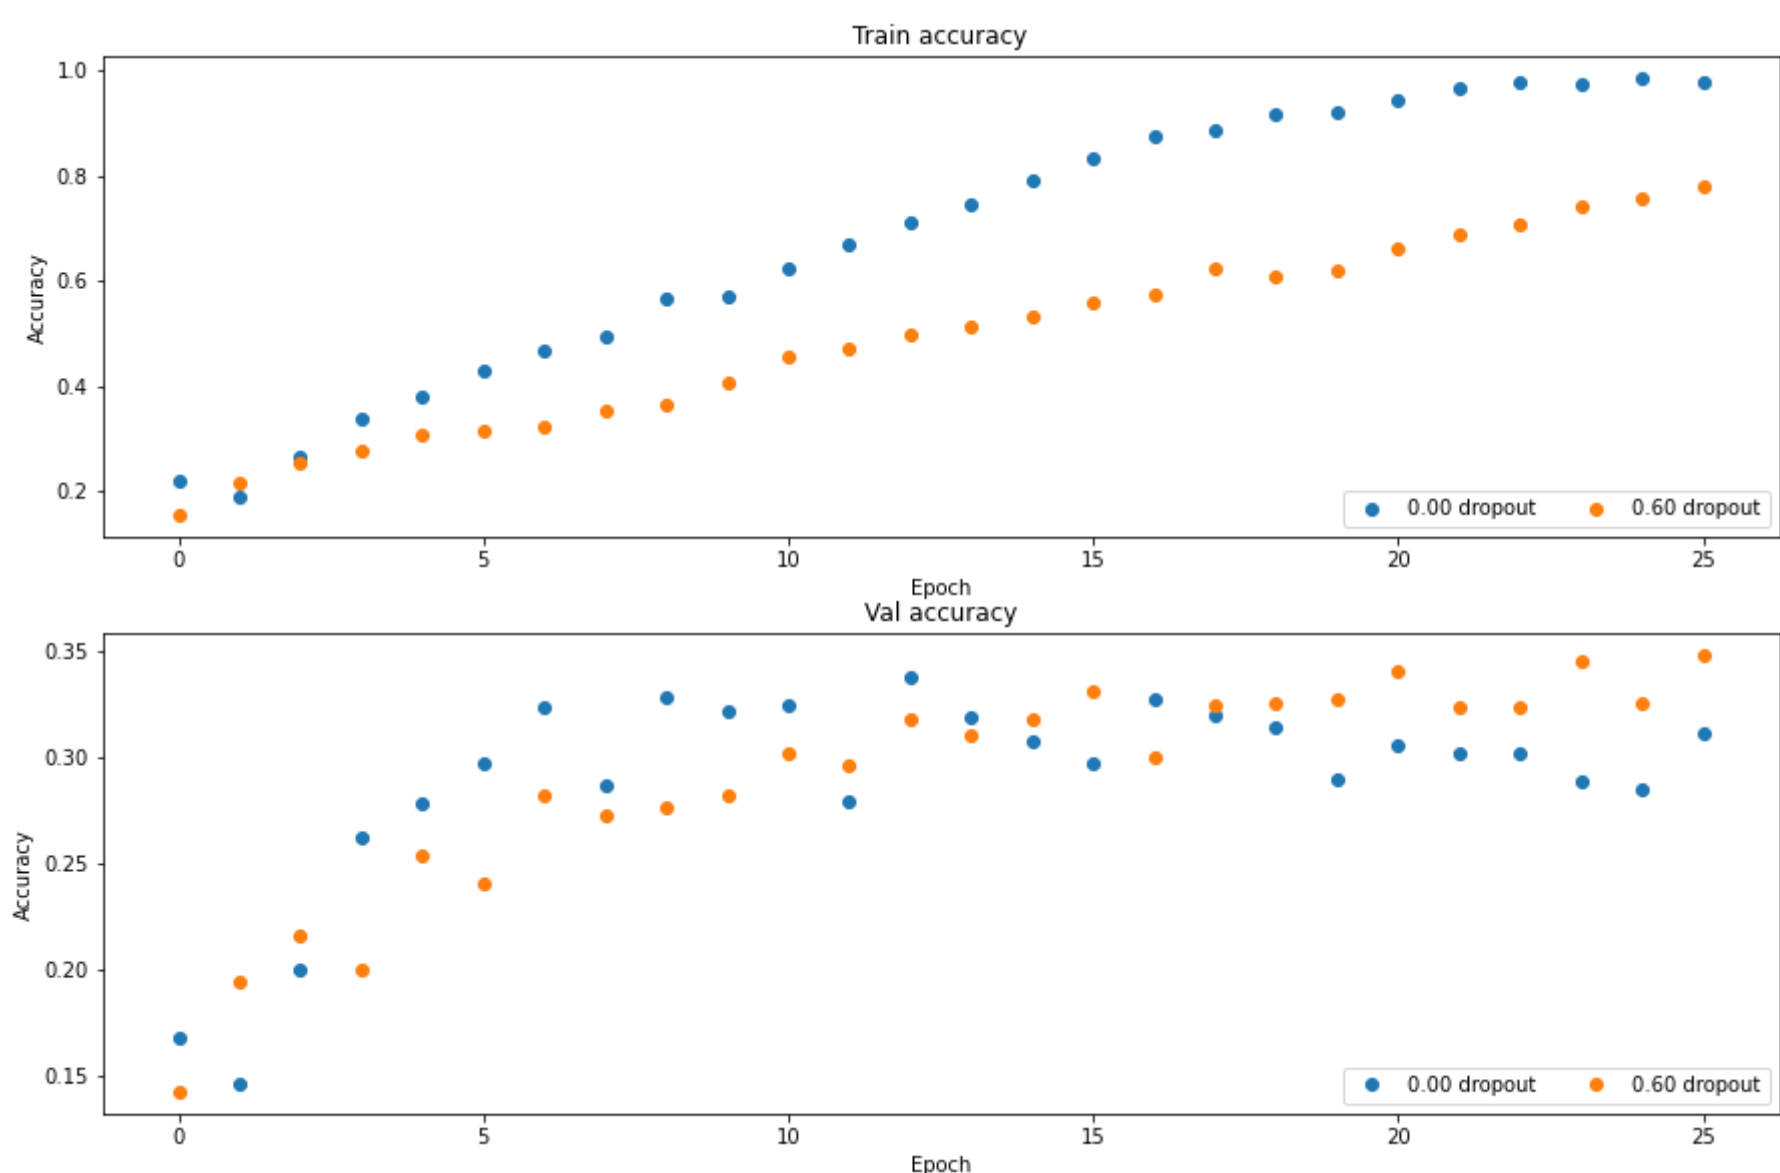
```

```

(Iteration 1 / 125) loss: 2.300804
(Epoch 0 / 25) train acc: 0.220000; val_acc: 0.168000
(Epoch 1 / 25) train acc: 0.188000; val_acc: 0.147000
(Epoch 2 / 25) train acc: 0.266000; val_acc: 0.200000
(Epoch 3 / 25) train acc: 0.338000; val_acc: 0.262000
(Epoch 4 / 25) train acc: 0.378000; val_acc: 0.278000
(Epoch 5 / 25) train acc: 0.428000; val_acc: 0.297000
(Epoch 6 / 25) train acc: 0.468000; val_acc: 0.323000
(Epoch 7 / 25) train acc: 0.494000; val_acc: 0.287000
(Epoch 8 / 25) train acc: 0.566000; val_acc: 0.328000
(Epoch 9 / 25) train acc: 0.572000; val_acc: 0.322000
(Epoch 10 / 25) train acc: 0.622000; val_acc: 0.324000
(Epoch 11 / 25) train acc: 0.670000; val_acc: 0.279000
(Epoch 12 / 25) train acc: 0.710000; val_acc: 0.338000
(Epoch 13 / 25) train acc: 0.746000; val_acc: 0.319000
(Epoch 14 / 25) train acc: 0.792000; val_acc: 0.307000
(Epoch 15 / 25) train acc: 0.834000; val_acc: 0.297000
(Epoch 16 / 25) train acc: 0.876000; val_acc: 0.327000
(Epoch 17 / 25) train acc: 0.886000; val_acc: 0.320000
(Epoch 18 / 25) train acc: 0.918000; val_acc: 0.314000
(Epoch 19 / 25) train acc: 0.922000; val_acc: 0.290000
(Epoch 20 / 25) train acc: 0.944000; val_acc: 0.306000
(Iteration 101 / 125) loss: 0.156105
(Epoch 21 / 25) train acc: 0.968000; val_acc: 0.302000
(Epoch 22 / 25) train acc: 0.978000; val_acc: 0.302000
(Epoch 23 / 25) train acc: 0.976000; val_acc: 0.289000
(Epoch 24 / 25) train acc: 0.986000; val_acc: 0.285000
(Epoch 25 / 25) train acc: 0.978000; val_acc: 0.311000
(Iteration 1 / 125) loss: 2.301328
(Epoch 0 / 25) train acc: 0.154000; val_acc: 0.143000
(Epoch 1 / 25) train acc: 0.214000; val_acc: 0.195000
(Epoch 2 / 25) train acc: 0.252000; val_acc: 0.216000
(Epoch 3 / 25) train acc: 0.276000; val_acc: 0.200000
(Epoch 4 / 25) train acc: 0.308000; val_acc: 0.254000
(Epoch 5 / 25) train acc: 0.316000; val_acc: 0.241000
(Epoch 6 / 25) train acc: 0.322000; val_acc: 0.282000
(Epoch 7 / 25) train acc: 0.354000; val_acc: 0.273000
(Epoch 8 / 25) train acc: 0.364000; val_acc: 0.276000
(Epoch 9 / 25) train acc: 0.408000; val_acc: 0.282000
(Epoch 10 / 25) train acc: 0.454000; val_acc: 0.302000
(Epoch 11 / 25) train acc: 0.472000; val_acc: 0.296000
(Epoch 12 / 25) train acc: 0.496000; val_acc: 0.318000
(Epoch 13 / 25) train acc: 0.512000; val_acc: 0.310000
(Epoch 14 / 25) train acc: 0.532000; val_acc: 0.318000
(Epoch 15 / 25) train acc: 0.558000; val_acc: 0.331000
(Epoch 16 / 25) train acc: 0.574000; val_acc: 0.300000
(Epoch 17 / 25) train acc: 0.624000; val_acc: 0.324000
(Epoch 18 / 25) train acc: 0.610000; val_acc: 0.325000
(Epoch 19 / 25) train acc: 0.620000; val_acc: 0.327000
(Epoch 20 / 25) train acc: 0.662000; val_acc: 0.340000
(Iteration 101 / 125) loss: 1.296187
(Epoch 21 / 25) train acc: 0.688000; val_acc: 0.323000
(Epoch 22 / 25) train acc: 0.708000; val_acc: 0.323000
(Epoch 23 / 25) train acc: 0.742000; val_acc: 0.345000
(Epoch 24 / 25) train acc: 0.756000; val_acc: 0.325000
(Epoch 25 / 25) train acc: 0.782000; val_acc: 0.348000

```

```
In [8]: 1 # Plot train and validation accuracies of the two models
2
3 train_accs = []
4 val_accs = []
5 for dropout in dropout_choices:
6     solver = solvers[dropout]
7     train_accs.append(solver.train_acc_history[-1])
8     val_accs.append(solver.val_acc_history[-1])
9
10 plt.subplot(3, 1, 1)
11 for dropout in dropout_choices:
12     plt.plot(solvers[dropout].train_acc_history, 'o', label='%.2f dropout' % dropout)
13 plt.title('Train accuracy')
14 plt.xlabel('Epoch')
15 plt.ylabel('Accuracy')
16 plt.legend(ncol=2, loc='lower right')
17
18 plt.subplot(3, 1, 2)
19 for dropout in dropout_choices:
20     plt.plot(solvers[dropout].val_acc_history, 'o', label='%.2f dropout' % dropout)
21 plt.title('Val accuracy')
22 plt.xlabel('Epoch')
23 plt.ylabel('Accuracy')
24 plt.legend(ncol=2, loc='lower right')
25
26 plt.gcf().set_size_inches(15, 15)
27 plt.show()
```



Question

Based off the results of this experiment, is dropout performing regularization? Explain your answer.

Answer:

By examining the plots of the accuracy for both validation and training data, we can conclude that dropout is performing regularization. If we look at the graph when dropout is not used, we can see that the final training accuracy is around 1 while the validation accuracy is around 0.3. This clearly indicates that the model has overfit to the training data. On the other hand, we can see that when dropout is used, the final training accuracy is around 0.8 while the final validation accuracy is around 0.35. From here, we can infer that when dropout is utilized, the discrepancy between the performance on training and validation data is lessened, indicating that the model is less likely to overfit. So, we can conclude that dropout serves as a method of regularization.

Final part of the assignment

Get over 55% validation accuracy on CIFAR-10 by using the layers you have implemented. You will be graded according to the following equation:

$$\min(\text{floor}((X - 32\%)) / 28\%, 1)$$

where if you get 60% or higher validation accuracy, you get full points.

In [19]:

```

1  # ===== #
2  # YOUR CODE HERE:
3  #   Implement a FC-net that achieves at least 55% validation accuracy
4  #   on CIFAR-10.
5  # ===== #
6  hiddenDims = [500,500,500]
7  weight_scale = 4e-2
8  dropout = 0.8
9  epochs = 17
10 batchSize = 490
11 optimizer = "adam"
12 lr = 5e-3
13 regularization = 0
14 lr_dec = 0.9
15
16 net = FullyConnectedNet(hiddenDims, weight_scale=weight_scale, dropout=dropout, use_batchnorm=True,
17                           reg = regularization)
18 solver = Solver(net, data,
19                 num_epochs=epochs, batch_size=batchSize,
20                 update_rule=optimizer,
21                 optim_config={
22                     'learning_rate': lr,
23                 },
24                 lr_decay = lr_dec, verbose=True, print_every=200)
25
26 solver.train()
27
28 # ===== #
29 # END YOUR CODE HERE
30 # ===== #

```

```

(Iteration 1 / 1700) loss: 2.448994
(Epoch 0 / 17) train acc: 0.145000; val_acc: 0.165000
(Epoch 1 / 17) train acc: 0.432000; val_acc: 0.440000
(Epoch 2 / 17) train acc: 0.509000; val_acc: 0.501000
(Iteration 201 / 1700) loss: 1.397535
(Epoch 3 / 17) train acc: 0.555000; val_acc: 0.511000
(Epoch 4 / 17) train acc: 0.595000; val_acc: 0.538000
(Iteration 401 / 1700) loss: 1.247360
(Epoch 5 / 17) train acc: 0.635000; val_acc: 0.540000
(Epoch 6 / 17) train acc: 0.650000; val_acc: 0.563000
(Iteration 601 / 1700) loss: 1.080137
(Epoch 7 / 17) train acc: 0.655000; val_acc: 0.560000
(Epoch 8 / 17) train acc: 0.690000; val_acc: 0.572000
(Iteration 801 / 1700) loss: 1.024190
(Epoch 9 / 17) train acc: 0.701000; val_acc: 0.561000
(Epoch 10 / 17) train acc: 0.719000; val_acc: 0.587000
(Iteration 1001 / 1700) loss: 0.967275
(Epoch 11 / 17) train acc: 0.769000; val_acc: 0.603000
(Epoch 12 / 17) train acc: 0.762000; val_acc: 0.603000
(Iteration 1201 / 1700) loss: 0.823831
(Epoch 13 / 17) train acc: 0.777000; val_acc: 0.588000
(Epoch 14 / 17) train acc: 0.809000; val_acc: 0.609000
(Iteration 1401 / 1700) loss: 0.694452
(Epoch 15 / 17) train acc: 0.843000; val_acc: 0.611000
(Epoch 16 / 17) train acc: 0.852000; val_acc: 0.601000
(Iteration 1601 / 1700) loss: 0.638089
(Epoch 17 / 17) train acc: 0.855000; val_acc: 0.603000

```



```

In [ ]: 1 import random
2
3 hiddenDimChoices = [100,200,300,400]
4 hiddenLayerAmounts = [1,2,3]
5 weight_scale = 2e-2
6 dropouts = [0.6,0.7,0.8,0.9]
7 epochs = [25]
8 batchSize = 490
9 optimizer = "adam"
10 lrs = [1e-3,5e-3,1e-2]
11 regularizations = [0,1e-8,1e-7]
12
13 valAcc = 0
14 while valAcc<0.6:
15     amountOfLayers = random.choice(hiddenLayerAmounts)
16     hiddenDims = []
17     for k in range(0,amountOfLayers):
18         hiddenDims.append(random.choice(hiddenDimChoices))
19     dropout = random.choice(dropouts)
20     epoch = random.choice(epochs)
21     lr = random.choice(lrs)
22     regularization = random.choice(regularizations)
23
24     net = FullyConnectedNet(hiddenDims, weight_scale=weight_scale, dropout=dropout, use_batchnorm=True,
25                             reg = regularization)
26     solver = Solver(net, data,
27                     num_epochs=epoch, batch_size=batchSize,
28                     update_rule=optimizer,
29                     optim_config={
30                         'learning_rate': lr,
31                     },
32                     verbose=True, print_every=200)
33
34     solver.train()
35     valAcc = solver.val_acc_history[-1]
36

```

```

(Iteration 1 / 2500) loss: 2.344163
(Epoch 0 / 25) train acc: 0.125000; val_acc: 0.135000
(Epoch 1 / 25) train acc: 0.435000; val_acc: 0.434000
(Epoch 2 / 25) train acc: 0.496000; val_acc: 0.492000
(Iteration 201 / 2500) loss: 1.486581
(Epoch 3 / 25) train acc: 0.510000; val_acc: 0.503000
(Epoch 4 / 25) train acc: 0.531000; val_acc: 0.514000
(Iteration 401 / 2500) loss: 1.374321
(Epoch 5 / 25) train acc: 0.555000; val_acc: 0.512000
(Epoch 6 / 25) train acc: 0.589000; val_acc: 0.516000
(Iteration 601 / 2500) loss: 1.272300
(Epoch 7 / 25) train acc: 0.581000; val_acc: 0.531000
(Epoch 8 / 25) train acc: 0.593000; val_acc: 0.536000
(Iteration 801 / 2500) loss: 1.186901
(Epoch 9 / 25) train acc: 0.642000; val_acc: 0.553000
(Epoch 10 / 25) train acc: 0.645000; val_acc: 0.550000
(Iteration 1001 / 2500) loss: 1.182109
(Epoch 11 / 25) train acc: 0.629000; val_acc: 0.545000
(Epoch 12 / 25) train acc: 0.650000; val_acc: 0.537000
(Iteration 1201 / 2500) loss: 1.044154

```

```

In [40]: 1 print(hiddenDims)
2 print(dropout)
3 print(epoch)
4 print(lr)
5 print(regularization)

```

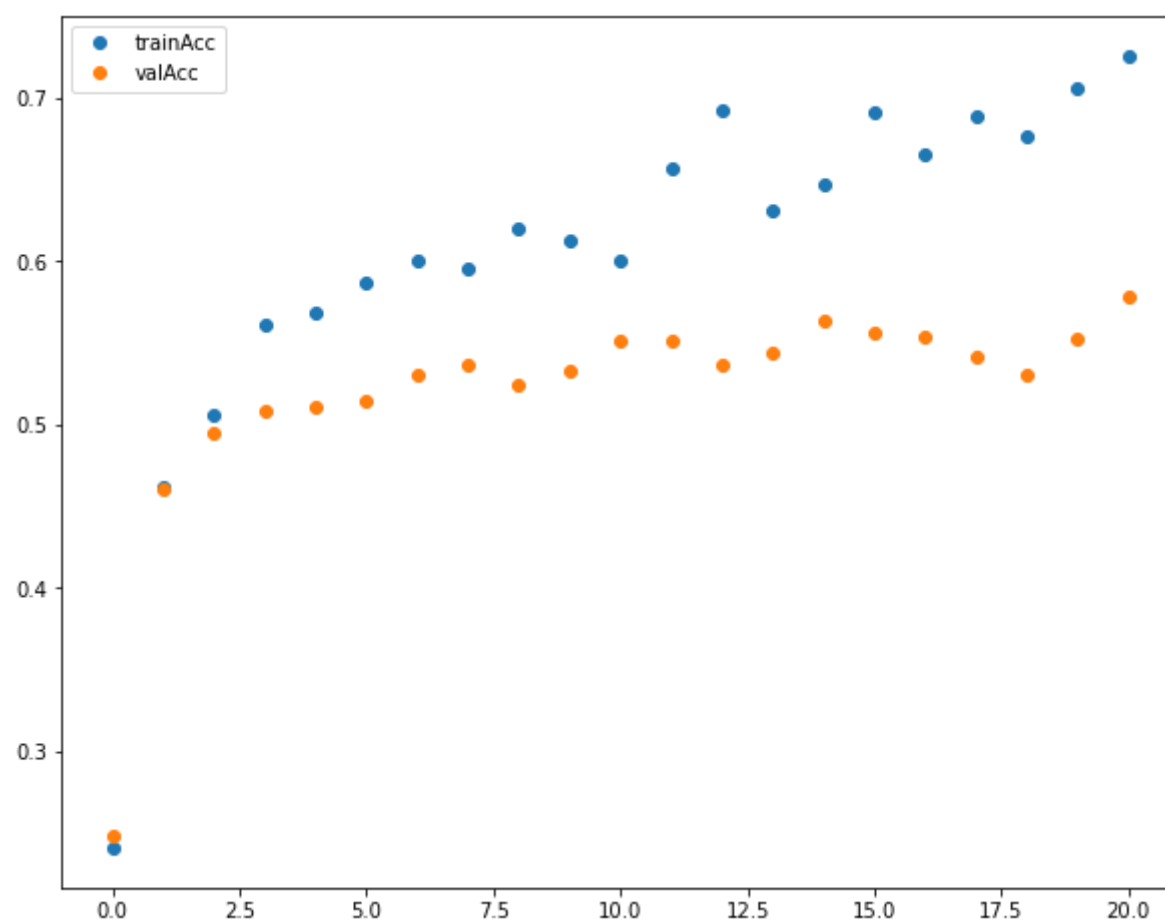
```

[200]
0.7
3
0.001
0

```

```
In [26]: 1 plt.figure()
2         plt.plot(solver.train_acc_history, 'o', label="trainAcc")
3         plt.plot(solver.val_acc_history, 'o', label="valAcc")
4         plt.legend()
```

Out[26]: <matplotlib.legend.Legend at 0x1ab242f3370>



```
In [31]: 1 solver.val_acc_history[-1]
```

Out[31]: 0.534