

```

import numpy as np
from .layers import *
from .layer_utils import *

class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network with ReLU nonlinearity and
    softmax loss that uses a modular layer design. We assume an input dimension
    of D, a hidden dimension of H, and perform classification over C classes.

    The architecture should be affine - relu - affine - softmax.

    Note that this class does not implement gradient descent; instead, it
    will interact with a separate Solver object that is responsible for running
    optimization.

    The learnable parameters of the model are stored in the dictionary
    self.params that maps parameter names to numpy arrays.
    """

    def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
                 dropout=0, weight_scale=1e-3, reg=0.0):
        """
        Initialize a new network.

        Inputs:
        - input_dim: An integer giving the size of the input
        - hidden_dims: An integer giving the size of the hidden layer
        - num_classes: An integer giving the number of classes to classify
        - dropout: Scalar between 0 and 1 giving dropout strength.
        - weight_scale: Scalar giving the standard deviation for random
          initialization of the weights.
        - reg: Scalar giving L2 regularization strength.
        """
        self.params = {}
        self.reg = reg

        # ===== #
        # YOUR CODE HERE:
        # Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
        # self.params['W2'], self.params['b1'] and self.params['b2']. The
        # biases are initialized to zero and the weights are initialized
        # so that each parameter has mean 0 and standard deviation weight_scale.
        # The dimensions of W1 should be (input_dim, hidden_dim) and the
        # dimensions of W2 should be (hidden_dims, num_classes)
        # ===== #

        # ===== #
        # END YOUR CODE HERE
        # ===== #

    def loss(self, X, y=None):
        """
        Compute loss and gradient for a minibatch of data.

        Inputs:
        - X: Array of input data of shape (N, d_1, ..., d_k)
        - y: Array of labels, of shape (N,). y[i] gives the label for X[i].

        Returns:
        If y is None, then run a test-time forward pass of the model and return:
        - scores: Array of shape (N, C) giving classification scores, where
          scores[i, c] is the classification score for X[i] and class c.

        If y is not None, then run a training-time forward and backward pass and
        return a tuple of:
        - loss: Scalar value giving the loss
        - grads: Dictionary with the same keys as self.params, mapping parameter
          names to gradients of the loss with respect to those parameters.
        """
        scores = None

        # ===== #
        # YOUR CODE HERE:
        # Implement the forward pass of the two-layer neural network. Store
        # the class scores as the variable 'scores'. Be sure to use the layers
        # you prior implemented.
        # ===== #

        # ===== #
        # END YOUR CODE HERE
        # ===== #

        # If y is None then we are in test mode so just return scores
        if y is None:
            return scores

        loss, grads = 0, {}
        # ===== #

```

```

# YOUR CODE HERE:
# Implement the backward pass of the two-layer neural net. Store
# the loss as the variable 'loss' and store the gradients in the
# 'grads' dictionary. For the grads dictionary, grads['W1'] holds
# the gradient for W1, grads['b1'] holds the gradient for b1, etc.
# i.e., grads[k] holds the gradient for self.params[k].
#
# Add L2 regularization, where there is an added cost 0.5*self.reg*W^2
# for each W. Be sure to include the 0.5 multiplying factor to
# match our implementation.
#
# And be sure to use the layers you prior implemented.
# ===== #

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

```

```
class FullyConnectedNet(object):
```

```

    """
    A fully-connected neural network with an arbitrary number of hidden layers,
    ReLU nonlinearities, and a softmax loss function. This will also implement
    dropout and batch normalization as options. For a network with L layers,
    the architecture will be

```

```
{affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
```

```

where batch normalization and dropout are optional, and the {...} block is
repeated L - 1 times.

```

```

Similar to the TwoLayerNet above, learnable parameters are stored in the
self.params dictionary and will be learned using the Solver class.
"""

```

```

def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
              dropout=0, use_batchnorm=False, reg=0.0,
              weight_scale=1e-2, dtype=np.float32, seed=None):

```

```

    """
    Initialize a new FullyConnectedNet.

```

```
Inputs:
```

- hidden\_dims: A list of integers giving the size of each hidden layer.
- input\_dim: An integer giving the size of the input.
- num\_classes: An integer giving the number of classes to classify.
- dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then the network should not use dropout at all.
- use\_batchnorm: Whether or not the network should use batch normalization.
- reg: Scalar giving L2 regularization strength.
- weight\_scale: Scalar giving the standard deviation for random initialization of the weights.
- dtype: A numpy datatype object; all computations will be performed using this datatype. float32 is faster but less accurate, so you should use float64 for numeric gradient checking.
- seed: If not None, then pass this random seed to the dropout layers. This will make the dropout layers deterministic so we can gradient check the model.

```

    """
    self.use_batchnorm = use_batchnorm
    self.use_dropout = dropout > 0
    self.reg = reg
    self.num_layers = 1 + len(hidden_dims)
    self.dtype = dtype
    self.params = {}

```

```

# ===== #
# YOUR CODE HERE:
# Initialize all parameters of the network in the self.params dictionary.
# The weights and biases of layer 1 are W1 and b1; and in general the
# weights and biases of layer i are Wi and bi. The
# biases are initialized to zero and the weights are initialized
# so that each parameter has mean 0 and standard deviation weight_scale.
#
# BATCHNORM: Initialize the gammas of each layer to 1 and the beta
# parameters to zero. The gamma and beta parameters for layer 1 should
# be self.params['gamma1'] and self.params['beta1']. For layer 2, they
# should be gamma2 and beta2, etc. Only use batchnorm if self.use_batchnorm
# is true and DO NOT do batch normalize the output scores.
# ===== #

```

```
for i in range(0, self.num_layers):
```

```

    name_W = 'W'+str(i+1)
    name_b = 'b'+str(i+1)
    name_gamma = 'gamma' + str(i+1)
    name_beta = 'beta' + str(i+1)

```

```

    if i == 0: #First
        self.params[name_W] = np.random.normal(loc=0.0, scale=weight_scale, size = (input_dim, hidden_dims[i]))
        self.params[name_b] = np.zeros(hidden_dims[i])

```

```

        if self.use_batchnorm:
            self.params[name_gamma] = np.ones(hidden_dims[i])
            self.params[name_beta] = np.zeros(hidden_dims[i])

    elif i == self.num_layers-1:          #Last
        self.params[name_W] = np.random.normal(loc=0.0, scale=weight_scale, size = (hidden_dims[i-1], num_classes))
        self.params[name_b] = np.zeros(num_classes)

    else:                                #Between
        self.params[name_W] = np.random.normal(loc=0.0, scale=weight_scale, size = (hidden_dims[i-1], hidden_dims[i]))
        self.params[name_b] = np.zeros(hidden_dims[i])
        if self.use_batchnorm:
            self.params[name_gamma] = np.ones(hidden_dims[i])
            self.params[name_beta] = np.zeros(hidden_dims[i])

# ===== #
# END YOUR CODE HERE
# ===== #

# When using dropout we need to pass a dropout_param dictionary to each
# dropout layer so that the layer knows the dropout probability and the mode
# (train / test). You can pass the same dropout_param to each dropout layer.
self.dropout_param = {}
if self.use_dropout:
    self.dropout_param = {'mode': 'train', 'p': dropout}
    if seed is not None:
        self.dropout_param['seed'] = seed

# With batch normalization we need to keep track of running means and
# variances, so we need to pass a special bn_param object to each batch
# normalization layer. You should pass self.bn_params[0] to the forward pass
# of the first batch normalization layer, self.bn_params[1] to the forward
# pass of the second batch normalization layer, etc.
self.bn_params = []
if self.use_batchnorm:
    self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]

# Cast all parameters to the correct datatype
for k, v in self.params.items():
    self.params[k] = v.astype(dtype)

def loss(self, X, y=None):
    """
    Compute loss and gradient for the fully-connected net.

    Input / output: Same as TwoLayerNet above.
    """
    X = X.astype(self.dtype)
    mode = 'test' if y is None else 'train'

    # Set train/test mode for batchnorm params and dropout param since they
    # behave differently during training and testing.
    if self.dropout_param is not None:
        self.dropout_param['mode'] = mode
    if self.use_batchnorm:
        for bn_param in self.bn_params:
            bn_param['mode'] = mode

    scores = None

    # ===== #
    # YOUR CODE HERE:
    # Implement the forward pass of the FC net and store the output
    # scores as the variable "scores".
    #
    # BATCHNORM: If self.use_batchnorm is true, insert a bathnorm layer
    # between the affine_forward and relu_forward layers. You may
    # also write an affine_batchnorm_relu() function in layer_utils.py.
    #
    # DROPOUT: If dropout is non-zero, insert a dropout layer after
    # every ReLU layer.
    # ===== #

    affineCaches = {}
    reluCaches = {}
    batchnormCaches = {}
    dropoutCaches = {}

    x = X

    for i in range(self.num_layers - 1):
        digit = str(i+1)

        x, affineCaches[digit] = affine_forward(x, self.params['W' + digit], self.params['b' + digit])

        if self.use_batchnorm:
            x, batchnormCaches[digit] = batchnorm_forward(x, self.params['gamma' + digit],
                                                         self.params['beta' + digit],
                                                         self.bn_params[i])

        x, reluCaches[digit] = relu_forward(x=x)

```

```

        if self.use_dropout:
            x, dropoutCaches[digit] = dropout_forward(x, self.dropout_param)

# Last layer
digit = str(self.num_layers)
scores, affineCaches[digit] = affine_forward(x, self.params['W' + digit],
                                              self.params['b' + digit])

# ===== #
# END YOUR CODE HERE
# ===== #

# If test mode return early
if mode == 'test':
    return scores

loss, grads = 0.0, {}
# ===== #
# YOUR CODE HERE:
# Implement the backwards pass of the FC net and store the gradients
# in the grads dict, so that grads[k] is the gradient of self.params[k]
# Be sure your L2 regularization includes a 0.5 factor.
#
# BATCHNORM: Incorporate the backward pass of the batchnorm.
#
# DROPOUT: Incorporate the backward pass of dropout.
# ===== #
loss, dz = softmax_loss(scores, y)

for i in range(self.num_layers, 0, -1):
    name_W = 'W' + str(i)
    name_b = 'b' + str(i)
    name_gamma = "gamma" + str(i)
    name_beta = "beta" + str(i)
    loss = loss + (0.5 * self.reg * np.sum(self.params[name_W]*self.params[name_W]))

    if i == self.num_layers:
        dh1, grads[name_W], grads[name_b] = affine_backward(dz, affineCaches[str(i)])
    else:
        if self.use_dropout:
            dh1 = dropout_backward(dh1, dropoutCaches[str(i)])
        dh1 = relu_backward(dh1, reluCaches[str(i)])
        if self.use_batchnorm:
            dh1, grads[name_gamma], grads[name_beta] = batchnorm_backward(dh1, batchnormCaches[str(i)])

        dh1, grads[name_W], grads[name_b] = affine_backward(dh1, affineCaches[str(i)])

    grads[name_W] = grads[name_W] + self.reg * self.params[name_W]

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

```