

```

import numpy as np

from nndl.layers import *
from nndl.conv_layers import *
from utils.fast_layers import *
from nndl.layer_utils import *
from nndl.conv_layer_utils import *

import pdb

class ThreeLayerConvNet(object):
    """
    A three-layer convolutional network with the following architecture:

    conv - relu - 2x2 max pool - affine - relu - affine - softmax

    The network operates on minibatches of data that have shape (N, C, H, W)
    consisting of N images, each with height H and width W and with C input
    channels.
    """

    def __init__(self, input_dim=(3, 32, 32), num_filters=32, filter_size=7,
                 hidden_dim=100, num_classes=10, weight_scale=1e-3, reg=0.0,
                 dtype=np.float32, use_batchnorm=False):
        """
        Initialize a new network.

        Inputs:
        - input_dim: Tuple (C, H, W) giving size of input data
        - num_filters: Number of filters to use in the convolutional layer
        - filter_size: Size of filters to use in the convolutional layer
        - hidden_dim: Number of units to use in the fully-connected hidden layer
        - num_classes: Number of scores to produce from the final affine layer.
        - weight_scale: Scalar giving standard deviation for random initialization
          of weights.
        - reg: Scalar giving L2 regularization strength
        - dtype: numpy datatype to use for computation.
        """
        self.use_batchnorm = use_batchnorm
        self.params = {}
        self.reg = reg
        self.dtype = dtype

        # ===== #
        # YOUR CODE HERE:
        # Initialize the weights and biases of a three layer CNN. To initialize:
        # - the biases should be initialized to zeros.
        # - the weights should be initialized to a matrix with entries
        #   drawn from a Gaussian distribution with zero mean and
        #   standard deviation given by weight_scale.
        # ===== #

        C, H, W = input_dim

        poolOutH = (H - 2) // 2 + 1
        poolOutW = (W - 2) // 2 + 1

        self.params["W1"] = np.random.normal(0, weight_scale, size=(num_filters, C, filter_size, filter_size))
        self.params["W2"] = np.random.normal(0, weight_scale, size=(num_filters * poolOutH * poolOutW, hidden_dim))
        self.params["W3"] = np.random.normal(0, weight_scale, size=(hidden_dim, num_classes))

        self.params["b1"] = np.zeros((num_filters,))
        self.params["b2"] = np.zeros((hidden_dim,))
        self.params["b3"] = np.zeros((num_classes,))

        # ===== #
        # END YOUR CODE HERE
        # ===== #

        for k, v in self.params.items():
            self.params[k] = v.astype(dtype)

    def loss(self, X, y=None):
        """
        Evaluate loss and gradient for the three-layer convolutional network.

        Input / output: Same API as TwoLayerNet in fc_net.py.
        """
        W1, b1 = self.params['W1'], self.params['b1']
        W2, b2 = self.params['W2'], self.params['b2']
        W3, b3 = self.params['W3'], self.params['b3']

        # pass conv_param to the forward pass for the convolutional layer
        filter_size = W1.shape[2]
        conv_param = {'stride': 1, 'pad': (filter_size - 1) // 2}

        # pass pool_param to the forward pass for the max-pooling layer
        pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

```

```

scores = None

# ===== #
# YOUR CODE HERE:
#   Implement the forward pass of the three layer CNN. Store the output
#   scores as the variable "scores".
# ===== #

out, conv_cache = conv_relu_pool_forward(X, W1, b1, conv_param, pool_param)
out, c1 = affine_relu_forward(out, W2, b2)
out, c2 = affine_forward(out, W3, b3)

scores = out

# ===== #
# END YOUR CODE HERE
# ===== #

if y is None:
    return scores

loss, grads = 0, {}
# ===== #
# YOUR CODE HERE:
#   Implement the backward pass of the three layer CNN. Store the grads
#   in the grads dictionary, exactly as before (i.e., the gradient of
#   self.params[k] will be grads[k]). Store the loss as "loss", and
#   don't forget to add regularization on ALL weight matrices.
# ===== #

loss, dout = softmax_loss(scores, y)

dout, dw, db = affine_backward(dout, c2)
grads["W3"] = dw + dw * self.reg
grads["b3"] = db

dout, dw, db = affine_relu_backward(dout, c1)
grads["W2"] = dw + dw * self.reg
grads["b2"] = db

dout, dw, db = conv_relu_pool_backward(dout, conv_cache)
grads["W1"] = dw + dw * self.reg
grads["b1"] = db

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

```

pass