

```

import numpy as np
from nndl.layers import *
import pdb

def conv_forward_naive(x, w, b, conv_param):
    """
    A naive implementation of the forward pass for a convolutional layer.

    The input consists of N data points, each with C channels, height H and width
    W. We convolve each input with F different filters, where each filter spans
    all C channels and has height HH and width HH.

    Input:
    - x: Input data of shape (N, C, H, W)
    - w: Filter weights of shape (F, C, HH, WW)
    - b: Biases, of shape (F,)
    - conv_param: A dictionary with the following keys:
        - 'stride': The number of pixels between adjacent receptive fields in the
            horizontal and vertical directions.
        - 'pad': The number of pixels that will be used to zero-pad the input.

    Returns a tuple of:
    - out: Output data, of shape (N, F, H', W') where H' and W' are given by
         $H' = 1 + (H + 2 * \text{pad} - \text{HH}) / \text{stride}$ 
         $W' = 1 + (W + 2 * \text{pad} - \text{WW}) / \text{stride}$ 
    - cache: (x, w, b, conv_param)
    """
    out = None
    pad = conv_param['pad']
    stride = conv_param['stride']

    # ===== #
    # YOUR CODE HERE:
    # Implement the forward pass of a convolutional neural network.
    # Store the output as 'out'.
    # Hint: to pad the array, you can use the function np.pad.
    # ===== #
    N, C, H, W = x.shape
    F, C, HH, WW = w.shape

    outputHeight = (H + 2 * pad - HH) // stride + 1
    outputWidth = (W + 2 * pad - WW) // stride + 1

    # output dimensions and initialization
    out = np.zeros((N, F, outputHeight, outputWidth))

    xPadded = np.pad(x, ((0,0), (0,0), (pad, pad), (pad, pad)), 'constant', constant_values=0)

    for i in np.arange(N): #For data points
        for j in np.arange(F): #For each filter
            for hIndex in np.arange(outputHeight):
                for wIndex in np.arange(outputWidth):
                    startH = hIndex*stride
                    startW = wIndex*stride
                    imagePart = xPadded[i, :, startH:(startH + HH), startW:(startW + WW)]
                    filterWeights = w[j, :, :, :]
                    filterBias = b[j]
                    out[i, j, hIndex, wIndex] = np.sum(imagePart*filterWeights) + filterBias

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    cache = (x, w, b, conv_param)
    return out, cache

def conv_backward_naive(dout, cache):
    """
    A naive implementation of the backward pass for a convolutional layer.

    Inputs:
    - dout: Upstream derivatives.
    - cache: A tuple of (x, w, b, conv_param) as in conv_forward_naive

    Returns a tuple of:
    - dx: Gradient with respect to x
    - dw: Gradient with respect to w
    - db: Gradient with respect to b
    """
    dx, dw, db = None, None, None

    N, F, out_height, out_width = dout.shape
    x, w, b, conv_param = cache

    stride, pad = [conv_param['stride'], conv_param['pad']]
    xpad = np.pad(x, ((0,0), (0,0), (pad,pad), (pad,pad)), mode='constant')

```

```

num_filts, _, f_height, f_width = w.shape

# ===== #
# YOUR CODE HERE:
# Implement the backward pass of a convolutional neural network.
# Calculate the gradients: dx, dw, and db.
# ===== #
db = np.zeros(b.shape)
dw = np.zeros(w.shape)
dx = np.zeros(x.shape)
dxdpad = np.zeros(xpad.shape)

for dataInd in range(N):
    for filterInd in range(F):
        db[filterInd] += np.sum(dout[dataInd, filterInd])
        for hInd in range(out_height):
            for wInd in range(out_width):
                startH = hInd*stride
                startW = wInd*stride
                dw[filterInd] += xpad[dataInd, :, startH:startH + f_height,
                                     startW:startW + f_width] * dout[dataInd, filterInd, hInd, wInd]
                dxdpad[dataInd, :, startH:startH + f_height,
                        startW:startW + f_width] += w[filterInd] * dout[dataInd, filterInd, hInd, wInd]

dx = dxdpad[:, :, pad:-pad, pad:-pad]
# ===== #
# END YOUR CODE HERE
# ===== #

return dx, dw, db

def max_pool_forward_naive(x, pool_param):
    """
    A naive implementation of the forward pass for a max pooling layer.

    Inputs:
    - x: Input data, of shape (N, C, H, W)
    - pool_param: dictionary with the following keys:
      - 'pool_height': The height of each pooling region
      - 'pool_width': The width of each pooling region
      - 'stride': The distance between adjacent pooling regions

    Returns a tuple of:
    - out: Output data
    - cache: (x, pool_param)
    """
    out = None

    # ===== #
    # YOUR CODE HERE:
    # Implement the max pooling forward pass.
    # ===== #
    poolH = pool_param['pool_height']
    poolW = pool_param['pool_width']
    stride = pool_param['stride']

    N, C, H, W = x.shape

    outH = (H - poolH) // stride + 1
    outW = (W - poolW) // stride + 1

    outShape = (N, C, outH, outW)
    out = np.zeros(outShape)

    for dataInd in range(N):
        for channelInd in range(C):
            for hInd in range(outH):
                for wInd in range(outW):
                    startH = hInd*stride
                    startW = wInd*stride
                    out[dataInd, channelInd, hInd, wInd] = np.max(x[dataInd, channelInd,
                                                                    startH:(startH + poolH),
                                                                    startW:(startW + poolW)])

    # ===== #
    # END YOUR CODE HERE
    # ===== #
    cache = (x, pool_param)
    return out, cache

def max_pool_backward_naive(dout, cache):
    """
    A naive implementation of the backward pass for a max pooling layer.

    Inputs:
    - dout: Upstream derivatives
    - cache: A tuple of (x, pool_param) as in the forward pass.

    Returns:
    - dx: Gradient with respect to x
    """

```

```

"""
dx = None
x, pool_param = cache
pool_height, pool_width, stride = pool_param['pool_height'], pool_param['pool_width'], pool_param['stride']

# ===== #
# YOUR CODE HERE:
# Implement the max pooling backward pass.
# ===== #
N, C, H, W = x.shape
_, _, outH, outW = dout.shape

dx = np.zeros(x.shape)

for dataInd in range(N):
    for channelInd in range(C):
        for hInd in range(outH):
            for wInd in range(outW):
                startH = hInd*stride
                startW = wInd*stride
                section = x[dataInd, channelInd, startH:startH + pool_height, startW:startW + pool_width]

                mask = section == np.max(section)

                dx[dataInd, channelInd, startH:startH + pool_height,
                    startW:startW + pool_width] += dout[dataInd, channelInd, hInd, wInd] * mask
# ===== #
# END YOUR CODE HERE
# ===== #

return dx

def spatial_batchnorm_forward(x, gamma, beta, bn_param):
    """
    Computes the forward pass for spatial batch normalization.

    Inputs:
    - x: Input data of shape (N, C, H, W)
    - gamma: Scale parameter, of shape (C,)
    - beta: Shift parameter, of shape (C,)
    - bn_param: Dictionary with the following keys:
      - mode: 'train' or 'test'; required
      - eps: Constant for numeric stability
      - momentum: Constant for running mean / variance. momentum=0 means that
        old information is discarded completely at every time step, while
        momentum=1 means that new information is never incorporated. The
        default of momentum=0.9 should work well in most situations.
      - running_mean: Array of shape (D,) giving running mean of features
      - running_var Array of shape (D,) giving running variance of features

    Returns a tuple of:
    - out: Output data, of shape (N, C, H, W)
    - cache: Values needed for the backward pass
    """
    out, cache = None, None

    # ===== #
    # YOUR CODE HERE:
    # Implement the spatial batchnorm forward pass.
    #
    # You may find it useful to use the batchnorm forward pass you
    # implemented in HW #4.
    # ===== #

    N, C, H, W = x.shape
    x = x.reshape((N*H*W,C))

    out, cache = batchnorm_forward(x, gamma, beta, bn_param)
    out = out.T
    out = out.reshape(C,N,H,W)
    out = out.swapaxes(0,1)

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return out, cache

def spatial_batchnorm_backward(dout, cache):
    """
    Computes the backward pass for spatial batch normalization.

    Inputs:
    - dout: Upstream derivatives, of shape (N, C, H, W)
    - cache: Values from the forward pass

    Returns a tuple of:
    - dx: Gradient with respect to inputs, of shape (N, C, H, W)
    - dgamma: Gradient with respect to scale parameter, of shape (C,)
    - dbeta: Gradient with respect to shift parameter, of shape (C,)
    """

```

```

'''
dx, dgamma, dbeta = None, None, None

# ===== #
# YOUR CODE HERE:
#   Implement the spatial batchnorm backward pass.
#
#   You may find it useful to use the batchnorm forward pass you
#   implemented in HW #4.
# ===== #
N, C, H, W = dout.shape

dout = dout.swapaxes(0,1)
dout = dout.reshape(C,N*H*W)
dout = dout.T

dx, dgamma, dbeta = batchnorm_backward(dout, cache)

dx = dx.reshape((N,C,H,W))
dgamma = dgamma.reshape((C,))
dbeta = dbeta.reshape((C,))

# ===== #
# END YOUR CODE HERE
# ===== #

return dx, dgamma, dbeta

```